

Syntheon

RedrobRank Agentic Recruiter OS

Deterministic, Transparent, Offline-First AI Ranking
Built for the India.Runs & Redrob Challenge

The Problem with LLM Parsing

Why traditional GenAI falls short in high-volume recruiting

Non-Deterministic & Hallucinatory

LLMs frequently invent justifications or fail to reference actual evidence from the candidate's history. You cannot blindly trust black-box rankings for hiring.

Slow & Expensive

Parsing 100k+ candidates through an LLM API takes days and incurs massive token costs. It's unscalable for an agency like Redrob receiving massive volume.

AI Talent Discovery Is Not Keyword Search

Understanding the real challenges in parsing AI profiles

1. Keyword stuffing

AI terms appear without real engineering depth.

2. Ghost candidates

Profiles may be inactive, unavailable, or low-response.

3. Domain mismatch

Support, accounting, HR, content, or operations roles can contain AI language.

4. Hallucinated reasons

A ranking system loses trust when the reason does not match candidate facts.


Intelligent Candidate Discovery

Our offline-first, deterministic ranking architecture

Core Pillars:

1. **Deterministic Multi-Signal Ranking:** Zero LLM API calls during ranking.
2. **Auditable Outputs:** Every recommendation references actual parsed facts.
3. **Local-First:** CPU-only, highly reproducible, runs entirely on the host.
4. **Risk Detection:** Instantly flags honeypots, keyword stuffers, and role mismatches.
5. **Massive Scale:** Streams and scores 100,000+ candidate JSONL seamlessly.

From 100K Profiles to a Grounded Top-100



100K processed

100 final rows

CPU-only

No LLM/API

Validator passed

What the Senior AI Engineer JD Needs

✓ Must-have signals

- Retrieval / vector search / semantic search
- Ranking / recommendation systems
- LLM / RAG / fine-tuning / NLP
- Production ML / MLOps
- Evaluation: NDCG, MRR, A/B testing
- Python + data/systems infrastructure

■ Risk signals

- Non-technical role mismatch
- Prompt-only or keyword-stuffed profiles
- Manager-only with weak technical depth
- Pure research without production evidence
- Weak availability or response behavior

Deterministic Multi-Signal Ranking

Scoring Stack:

1. TF-IDF lexical recall
2. Structured AI skill score
3. Skill trust: proficiency, duration, endorsements
4. Career relevance with recency weighting
5. Production + evaluation evidence
6. Behavioral availability
7. Risk penalties
8. RRF fusion + candidate_id tie-break

Every Recommendation Is Auditable

Reasoning uses actual:

- Current title & Location
- YOE
- Matched AI/retrieval terms
- Production/eval evidence
- Redrob behavioral signals

Metrics & Proof:

- Official validator: Passed
- Internal validator: Passed
- Forensic audit: Passed
- Hallucination failures: 0
- Unique scores: 74
- Unique reasonings: 100

Local-First, Privacy-First Architecture

CLI Ranker:

rank.py, stream reader, feature extractor, scorer, CSV writer, reports

Audit Layer:

validator, forensic audit, security scans, quality reports

Interfaces:

React/Vite dashboard, Streamlit sample sandbox

No external APIs

No keys

No LLM calls

No full dataset browser upload

Final Output Surfaces Real AI Profiles

100,000 processed | 100 ranked | 74 unique scores | 100 unique reasonings
0 hallucination failures | CSV validator passed | Under 5 MB output

CAND_0018499	Senior ML Engineer	7.2 YOE	Noida
CAND_0005260	Senior NLP Engineer	5.2 YOE	Chennai
CAND_0046525	Senior ML Engineer	6.1 YOE	Pune
CAND_0081846	Lead AI Engineer	6.7 YOE	Jaipur
CAND_0042506	Search Engineer	4.2 YOE	Mumbai

Final Submission Package

Team ID: Syntheon | **Team Name:** Syntheon

GitHub Repo: <https://github.com/Sauravssoni/Screener>

Ranked CSV: submissions/submission.csv

Deck PDF: docs/RedrobRank_Official_Submission_Deck.pdf

Audit: reports/final_ranking_audit.md

Dashboard Demo: npm run dev | **Sandbox App:** streamlit run sandbox_app.py

**RedrobRank is a reproducible, auditable TalentOps engine —
not just a ranking script.**